
Real-Time Systems

Last modified: 02.06.2021

by L.J. Opalski

Overview of Real-Time Systems

- A real-time (RT) system requires that results be produced within a specified deadline period.
 - An embedded system is a computing device that is part of a larger system (i.e. automobile, airliner)
 - A safety-critical system is a real-time system with catastrophic results in case of failure
- A hard real-time system guarantees that real-time tasks be completed within their required deadlines
- A soft real-time system provides priority of real-time tasks over non real-time tasks


Features of Real-Time Kernels

- Most real-time systems do not provide the features found in a standard desktop system. Reasons include:
 - Real-time systems are typically single-purpose
 - Real-time systems often do not require interfacing with a user
 - Features found in a desktop PC require more substantial hardware than what is typically available in a real-time system
 - Cost-effective features may introduce non-deterministic latency
 - Real-time applications must be able to control the execution sequence of their concurrent process in order to meet externally defined response requirements.

Memory in Real-Time Systems

- Address translation may occur via:


- (1) Real-addressing mode where programs generate actual addresses
- (2) Relocation register mode
- (3) Implementing full virtual memory



- The concept of **memory locking** is introduced in RT systems to eliminate the indeterminacy introduced by paging and swapping, and to support an **upper bound** on the time required to access the memory mapped into the address space of a process.


“Must have” for Real-Time Systems

- In general, real-time operating systems must provide:
 - (1) **Preemptive**, priority-based, **scheduling**
 - (2) **Preemptive kernels**
 - (3) Minimized **event latency**




Interrupt Latency

- **Interrupt latency** is the period of time from when an interrupt arrives at the CPU to when it is serviced
- In RT systems interrupt latency must be **bounded** to guarantee the **deterministic behavior** of the system.




Dispatch Latency

- Dispatch latency is the amount of time required for the scheduler to stop one process and start another
- Dispatch latency can be reduced by:
 - preemption points in long-duration system calls
 - preemptive kernel



Real-Time CPU Scheduling

- Important characteristics of tasks, which RT processes realize:
 - Aperiodic task – has to be started or finished by specified time
 - Periodic task – require the CPU at specified (periodic) intervals
 - p is the duration of the period
 - d is the deadline by when the process must be serviced
 - t is the processing time



- Necessary condition of scheduling feasibility: $\sum_{i=1}^N \frac{t_i}{p_i} \leq 1$

Scheduling of tasks when P_2 has a higher priority than P_1

$p_1=50, p_2=100$ – periods

$t_1=20, t_2=35$ – processing times


CPU utilization:

$$P_1: t_1/p_1 = 20/50 = 0.4$$

$$P_2: t_2/p_2 = 35/100 = 0.35$$

Both: $0.4 + 0.35 = 0.75 < 1 \rightarrow$ it seems that we might meet deadlines


Surprise (?)



Priority-based CPU allocation **failed** in this example for assumed priorities

Rate Monotonic Scheduling (RMS)

- Rate Monotonic Scheduling (RMS) algorithm assumes that a process priority is assigned based on the inverse of its period
- Shorter periods = higher priority;
- Longer periods = lower priority
- P_1 is assigned a higher priority than P_2 . ($t_1=20$; $t_2=35$)



- This time the priorities resulted in meeting deadlines. Is the RM scheduling always successful?

Missed Deadlines with Rate Monotonic Scheduling


- Rate-monotonic scheduling is considered optimal in that if a set of processes cannot be scheduled by this algorithm, it cannot be scheduled by any other algorithm that assigns static priorities.
- Example.

$p_1=50, p_2=80$ – periods; $t_1=25, t_2=35$ – processing times

CPU utilization:

$$P_1: t_1/p_1=25/50=0.5; \quad P_2: t_2/p_2=35/80=0.44$$

Necessary condition: $0.5+0.44=0.94<1$ is satisfied



Note:


- The worst-case CPU utilization for scheduling N processes is bounded by $N(2^{1/N}-1)$. For $N=2$ it makes $0.83 < 0.94$ - hence the above failure. $\lim_{N \rightarrow \infty} N(2^{1/N} - 1) = \ln(2) \approx 0.6931$
- RMS (a priority scheduler) can be affected by **priority inversion**

Earliest Deadline First (EDF) Scheduling

- **Earliest Deadline First (EDF)** Scheduling assigns priorities (dynamically) according to current **deadlines**:

the earlier the deadline, the higher the priority;

the later the deadline, the lower the priority



- Unlike the rate-monotonic algorithm EDF scheduling does not require that processes be periodic, nor must a process require a constant CPU time per burst.
- EDF only requires that a process announces its deadline to the scheduler when it becomes runnable.

POSIX Real-Time Scheduling

The POSIX.1b standard

- Each thread is controlled by an associated **scheduling policy** and **priority**.
- The scheduling **contention scope** of a thread defines the set of threads with which the thread competes for use of the processing resources. System or process scope.
- There is, conceptually, one thread list of runnable threads for each priority. A **scheduling policy** defines allowable operations on the set of lists (including moving threads between and within lists). Associated with each policy is a **priority range**.
- The thread, which is at the head of the highest priority non-empty thread list becomes a running thread, regardless of its associated policy.

POSIX Real-Time Thread Scheduling

Four thread scheduling policies are required by POSIX 1003.1 Standard:


- Scheduling classes for real-time threads:
 1. **SCHED_FIFO** - threads are scheduled using a FCFS strategy with a FIFO queue. There is no time-slicing for threads of equal priority
 2. **SCHED_RR** - similar to SCHED_FIFO except time-slicing occurs for threads of equal priority
 3. **SCHED_SPORADIC** – high priority for at most specified time, if the time is exceeded – the priority drops; the execution capacity is replenished after some specified time.
- Another scheduling policy
 4. **SCHED_OTHER** – can behave like FIFO or RR policy, but it is not a RT policy, so all RT threads preempt threads with SCHED_OTHER policy. Threads subject to this policy can have dynamic priorities (e.g. to implement ageing).

Linux: additional classes of no-RT processes: SCHED_BATCH, SCHED_IDLE and of RT processes: **SCHED_DEADLINE** (Global Earliest Deadline First scheduling) – instead of SCHED_SPORADIC.
See man sched(7).

POSIX RT profiles

- POSIX 1003.1 Standard
 - Allows writing portable real-time applications
 - Very large: inappropriate for embedded real-time systems
- POSIX 1003.13 – defines four real-time system subsets (profiles)
 - Minimal (PSE51): for small embedded systems. Platform: Small embedded system, with no MMU, no disk, no terminal. Model: controller of a “Toaster”
 - Controller (PSE52): for industrial controllers. Platform: Special purpose controller, with no MMU, but with a disk containing a simplified file system. Model: industrial robot controller
 - Dedicated (PSE53): for large embedded systems. Platform: Large embedded system with file system on disk, with an MMU; software is complex and requires memory protection and network communications. Models: avionics controller, cellular phone cell node.
 - Multi-Purpose (PSE54): for large general-purpose systems with RT requirements. Platform: Large real-time system with all the features, including a development environment, network communications, file system on disk, terminal and GUI, etc. Model: workstation with RT requirements: e.g. for air traffic control systems, or for telemetry systems for Formula One racing cars.

Example RT system: VxWorks



Architectures: ARM (64b, v7), Intel (IA), Power PC, RISC-V,....

Multiprocessor systems: AMP i SMP; multi-OS (Type 1 hypervisor)

Open source Raspberry Pi Board Support Package

<https://www.windriver.com/products/vxworks>

Wind Microkernel

■ The Wind microkernel features:

- small size (20kB)
- multitasking with processes and threads
- preemptive and non-preemptive round-robin scheduling
- manages interrupts (with bounded interrupt and dispatch latency times)
- shared memory and message passing interprocess communication support

■ There is a support to two levels of virtual memory

- Control of the page cache on a per-page basis, eg. allowing for non-cacheable pages
- Optional virtual memory component VxVMI along with hardware MMU allow selected data areas to be marked as private (accessible from a specific task)

■ POSIX PSE52 interface support.

■ Many certificates available, handy for development of aerospace, automobile, medical and other safety-critical systems

	DO-178C certifiable to DAL A
	ISO 26262 certifiable to ASIL-D
	IEC 61508 certifiable to SIL3
	IEC 62304 Class C Risk Level

RTLinux

- RTLinux is a hard real-time microkernel
- In RTLinux the standard Linux kernel runs as a fully preemptive task in a small operating system
- The real-time (RT) kernel handles all interrupts – directing each interrupt to a handler in the standard kernel or to an interrupt handler in the RT kernel.
- RTLinux prevents the standard Linux kernel from ever disabling interrupts, thus ensuring that it cannot add latency to the RT system
- RTLinux provides additional scheduling policies, including rate-monotonic scheduling (RMS) and earliest-deadline-first scheduling (EDF).